

Data-Intensive Computing with MapReduce

Session 2: Hadoop Nuts and Bolts

Jimmy Lin
University of Maryland
Thursday, January 31, 2013



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details



Source: Wikipedia (The Scream)



Source: Wikipedia (Japanese rock garden)

How will I actually learn Hadoop?

- This class session
- Hadoop: The Definitive Guide
- RTFM
- RTFC(!)



Basic Hadoop API*

- Mapper

- `void setup(Mapper.Context context)`
Called once at the beginning of the task
- `void map(K key, V value, Mapper.Context context)`
Called once for each key/value pair in the input split
- `void cleanup(Mapper.Context context)`
Called once at the end of the task

- Reducer/Combiner

- `void setup(Reducer.Context context)`
Called once at the start of the task
- `void reduce(K key, Iterable<V> values, Reducer.Context context)`
Called once for each key
- `void cleanup(Reducer.Context context)`
Called once at the end of the task

*Note that there are two versions of the API!

Basic Hadoop API*

- Partitioner

- `int getPartition(K key, V value, int numPartitions)`
Get the partition number given total number of partitions

- Job

- Represents a packaged Hadoop job for submission to cluster
- Need to specify input and output paths
- Need to specify input and output formats
- Need to specify mapper, reducer, combiner, partitioner classes
- Need to specify intermediate/final key/value classes
- Don't depend on defaults!

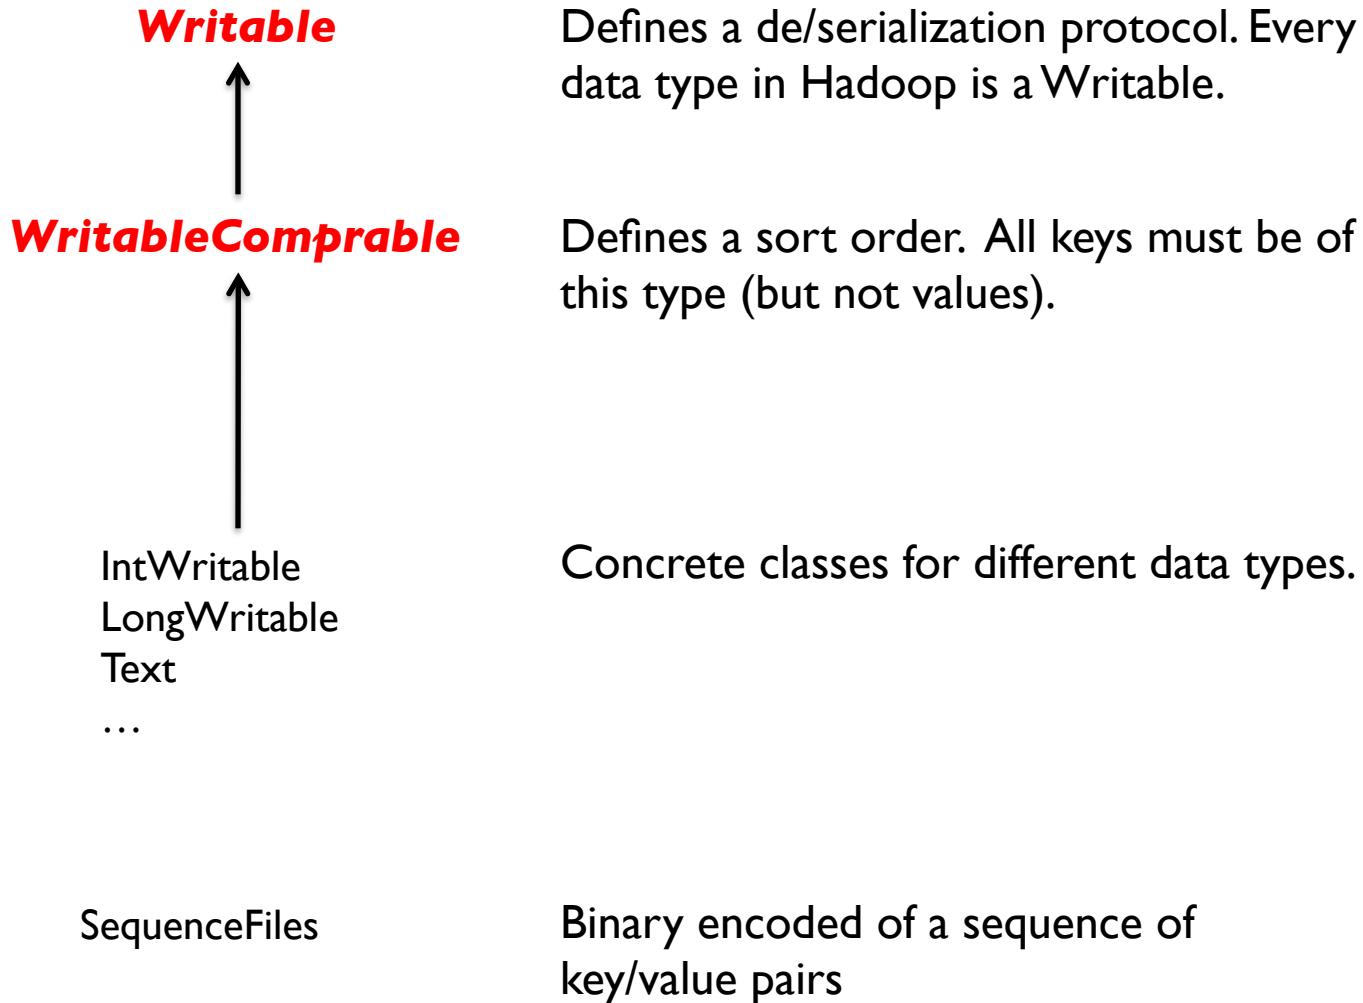
*Note that there are two versions of the API!

A tale of two packages...

org.apache.hadoop.mapreduce
org.apache.hadoop.mapred



Data Types in Hadoop: Keys and Values



“Hello World”: Word Count

Map(String docid, String text):

for each word w in text:

 Emit(w, 1);

Reduce(String term, Iterator<Int> values):

 int sum = 0;

 for each v in values:

 sum += v;

 Emit(term, value);

“Hello World”: Word Count

```
private static class MyMapper extends  
    Mapper<LongWritable, Text, Text, IntWritable> {  
  
    private final static IntWritable ONE = new IntWritable(1);  
    private final static Text WORD = new Text();  
  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
        String line = ((Text) value).toString();  
        StringTokenizer itr = new StringTokenizer(line);  
        while (itr.hasMoreTokens()) {  
            WORD.set(itr.nextToken());  
            context.write(WORD, ONE);  
        }  
    }  
}
```

“Hello World”: Word Count

```
private static class MyReducer extends  
    Reducer<Text, IntWritable, Text, IntWritable> {  
  
    private final static IntWritable SUM = new IntWritable();  
  
    @Override  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context) throws IOException, InterruptedException {  
        Iterator<IntWritable> iter = values.iterator();  
        int sum = 0;  
        while (iter.hasNext()) {  
            sum += iter.next().get();  
        }  
        SUM.set(sum);  
        context.write(key, SUM);  
    }  
}
```

Three Gotchas

- Avoid object creation at all costs
 - Reuse Writable objects, change the payload
- Execution framework reuses value object in reducer
- Passing parameters into mappers and reducers:
 - Directly in the Job object for parameters
 - DistributedCache for larger static data

Complex Data Types in Hadoop

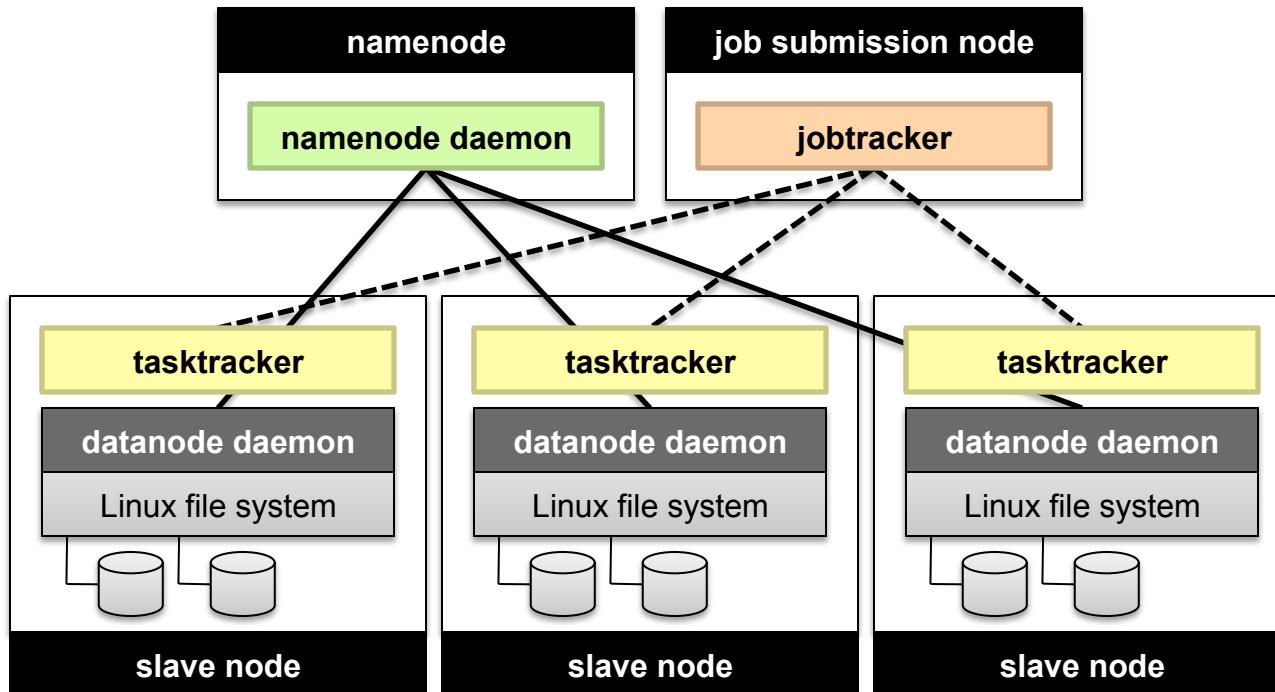
- How do you implement complex data types?
- The easiest way:
 - Encoded it as Text, e.g., (a, b) = “a:b”
 - Use regular expressions to parse and extract data
 - Works, but pretty hack-ish
- The hard way:
 - Define a custom implementation of Writable(Comparable)
 - Must implement: readFields, write, (compareTo)
 - Computationally efficient, but slow for rapid prototyping
 - Implement WritableComparator hook for performance
- Somewhere in the middle:
 - Cloud⁹ offers JSON support and lots of useful Hadoop types

Basic Cluster Components*

- One of each:
 - Namenode (NN): master node for HDFS
 - Jobtracker (JT): master node for job submission
- Set of each per slave machine:
 - Tasktracker (TT): contains multiple task slots
 - Datanode (DN): serves HDFS data blocks

* Not quite... leaving aside YARN for now

Putting everything together...



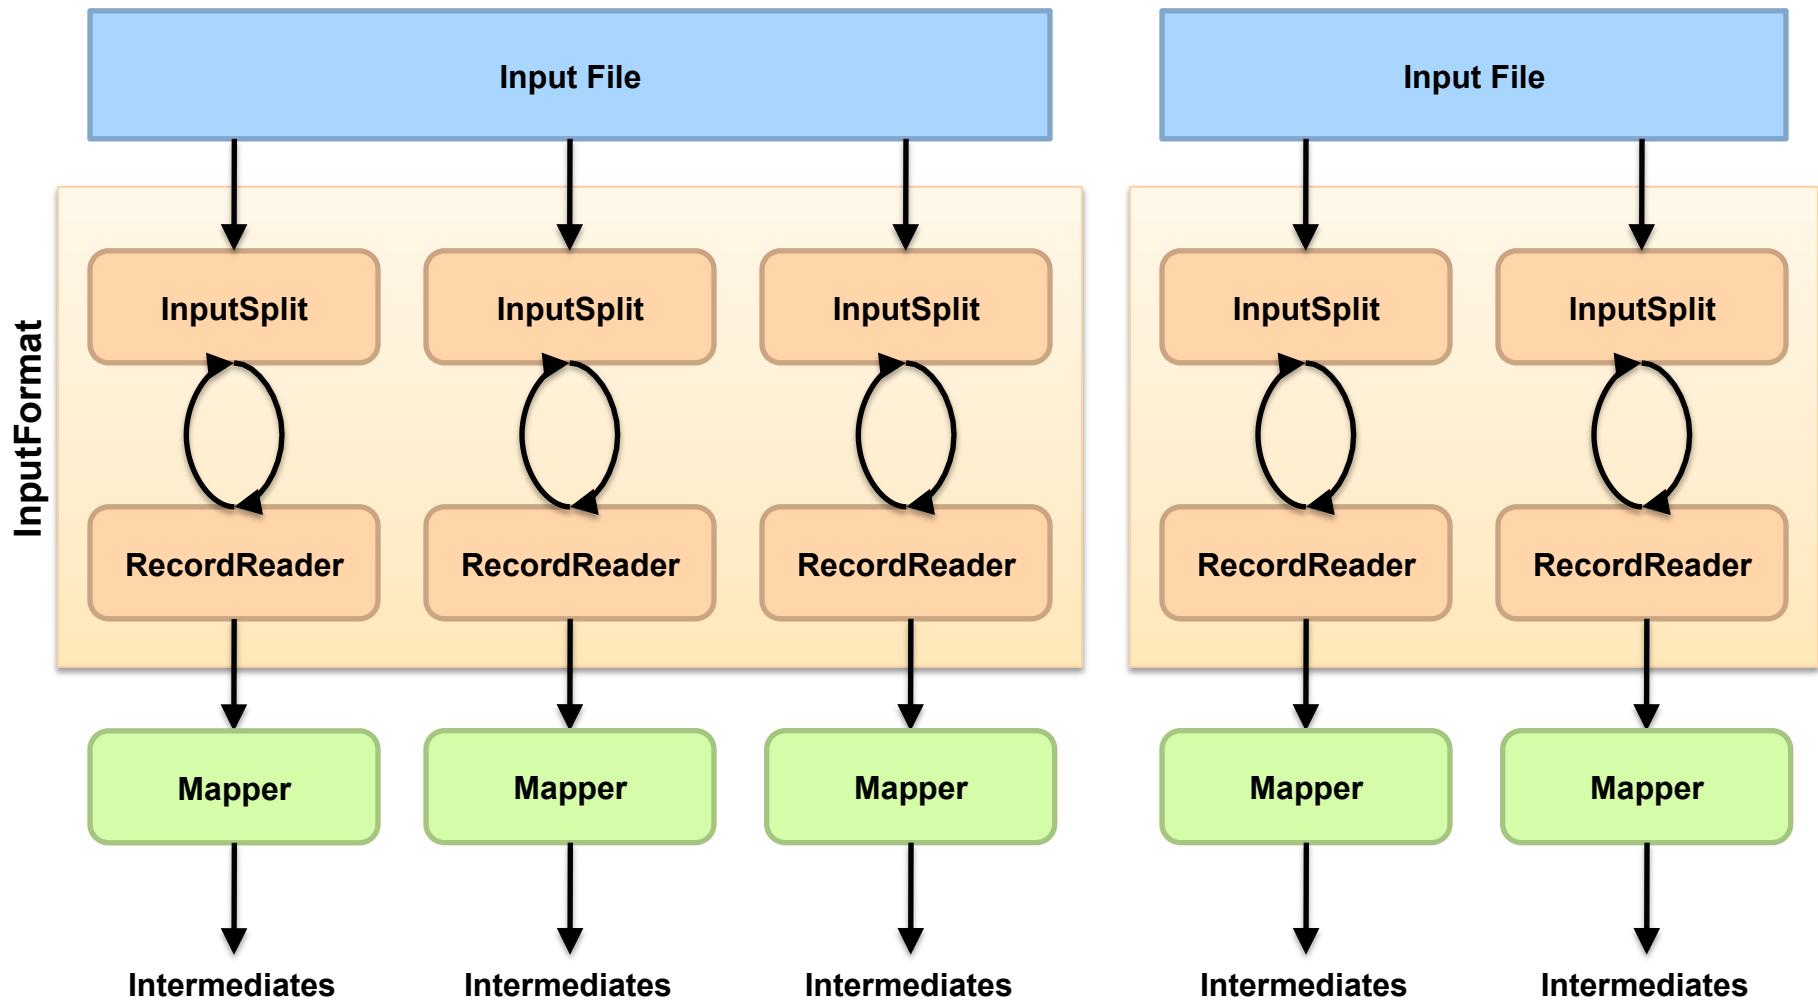
Anatomy of a Job

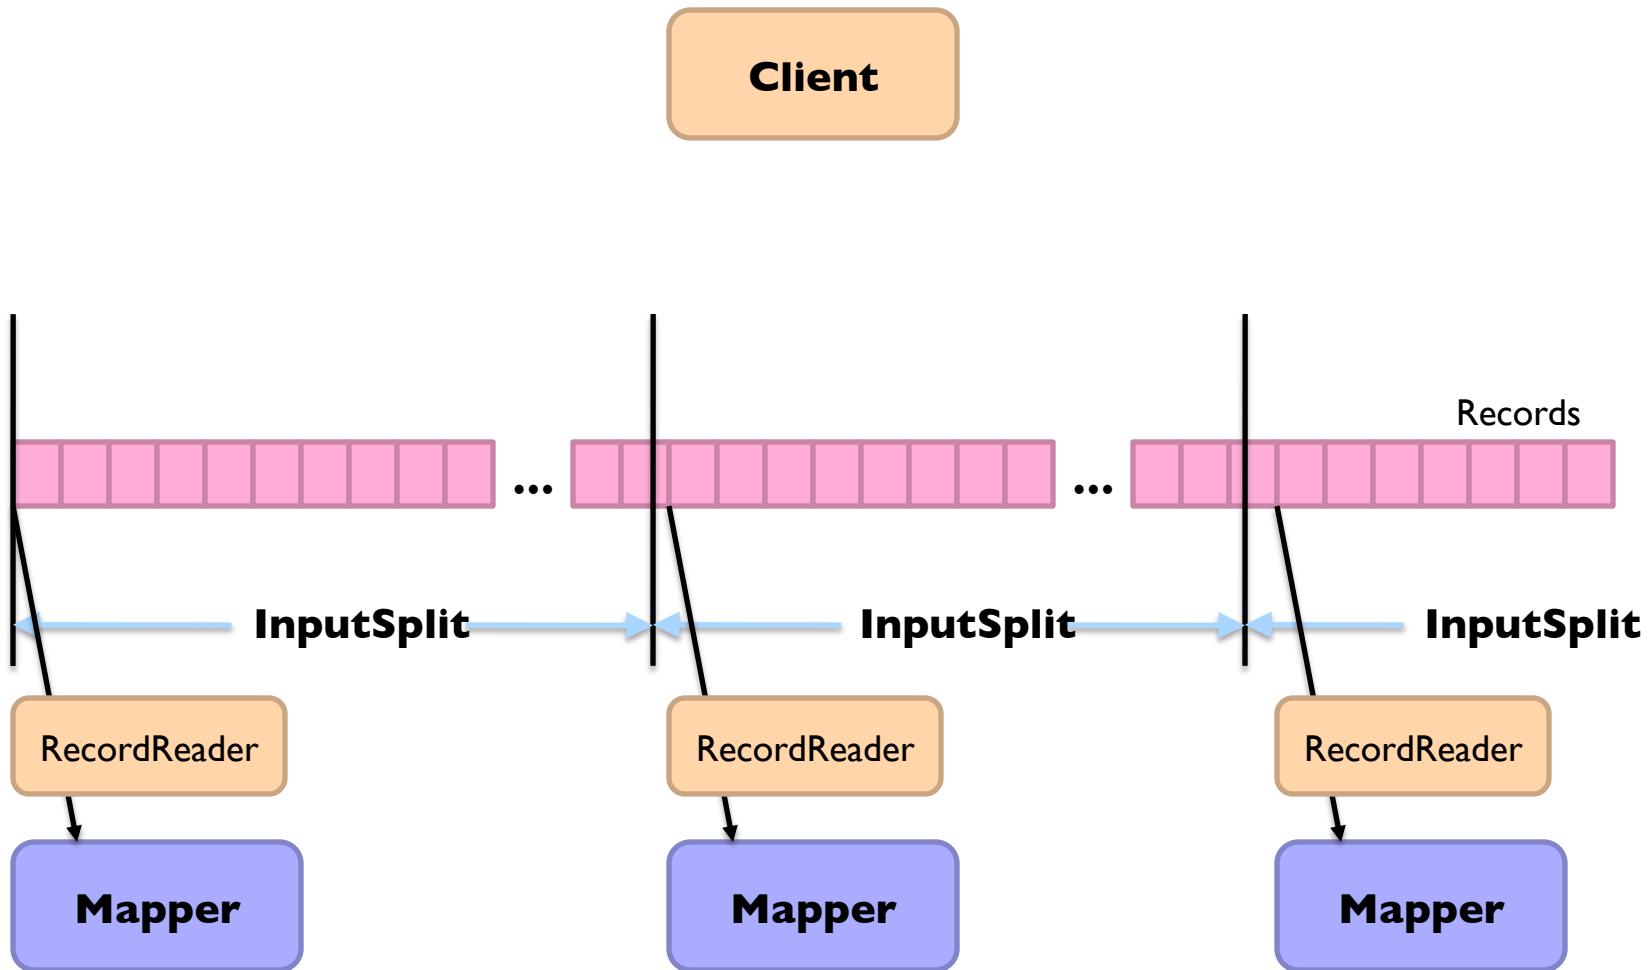
- MapReduce program in Hadoop = Hadoop job
 - Jobs are divided into map and reduce tasks
 - An instance of running a task is called a task attempt (occupies a slot)
 - Multiple jobs can be composed into a workflow
- Job submission:
 - Client (i.e., driver program) creates a job, configures it, and submits it to jobtracker
 - That's it! The Hadoop cluster takes over...

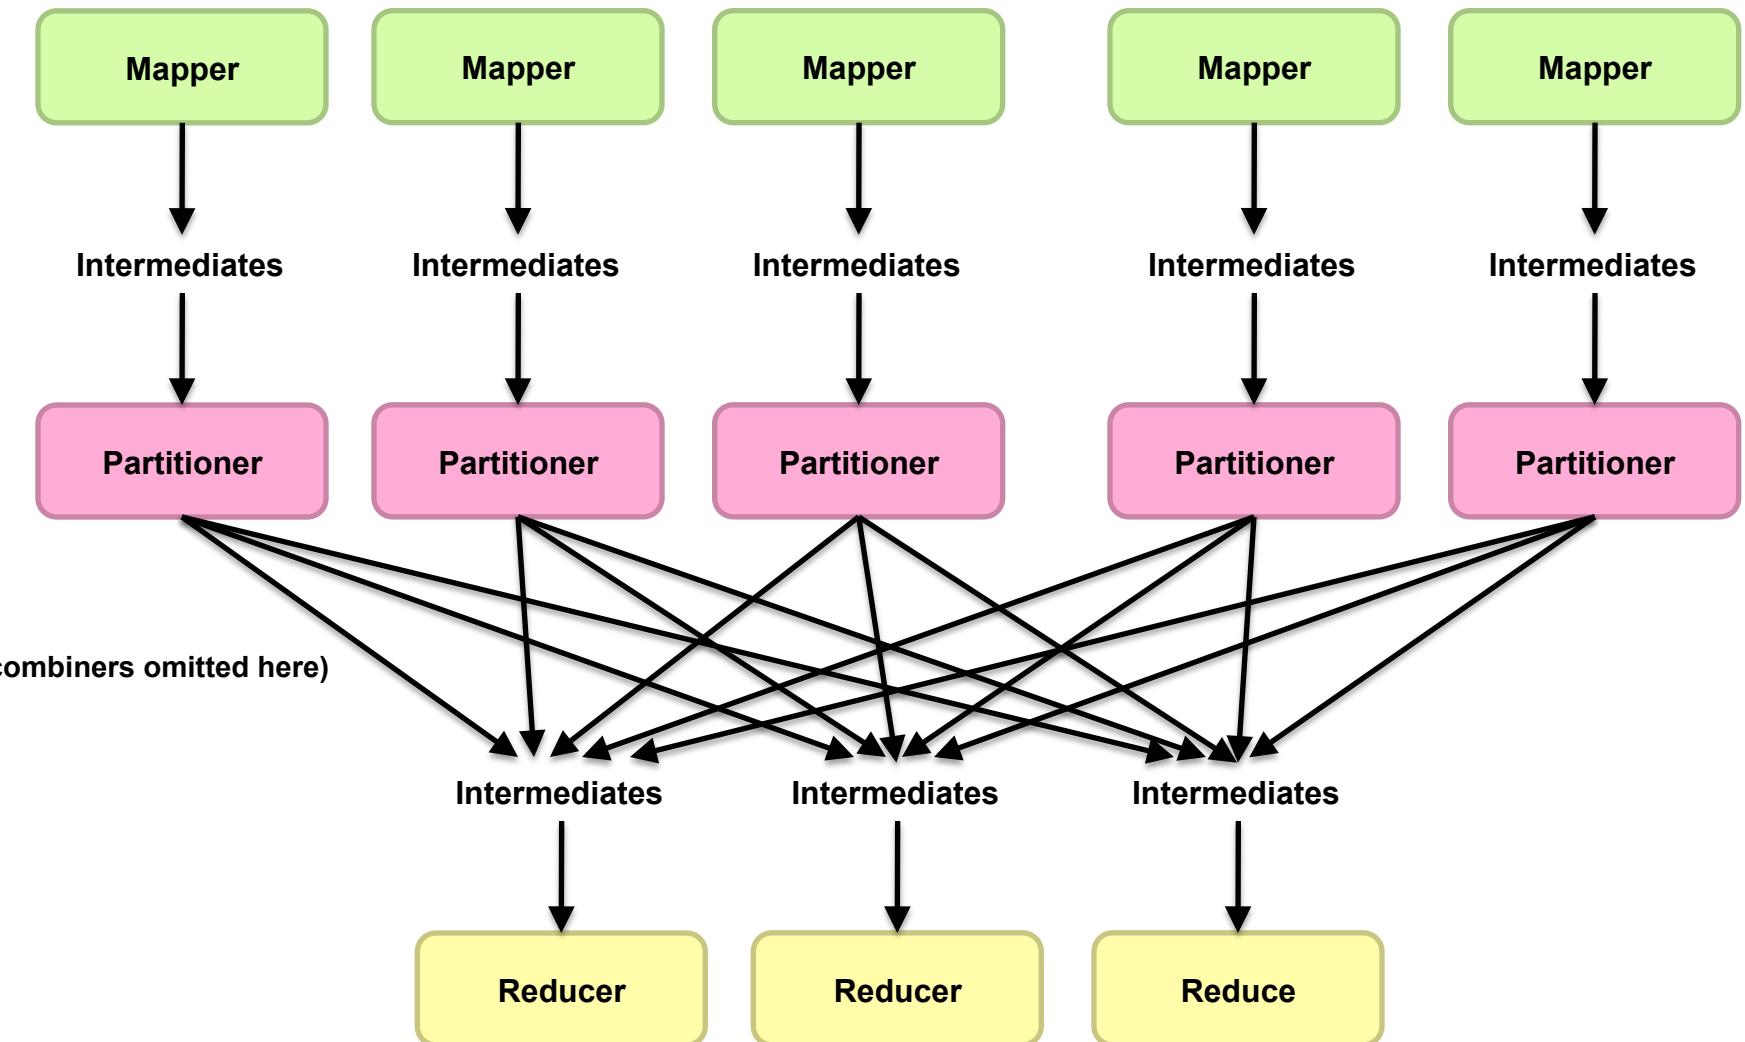
Anatomy of a Job

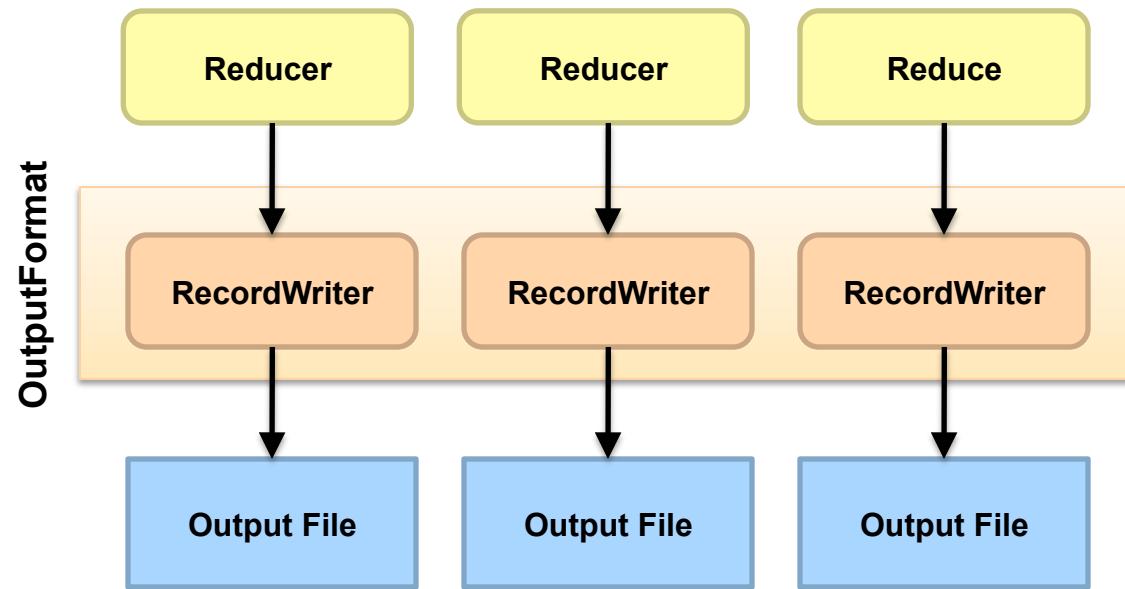
- Behind the scenes:

- Input splits are computed (on client end)
- Job data (jar, configuration XML) are sent to JobTracker
- JobTracker puts job data in shared location, enqueues tasks
- TaskTrackers poll for tasks
- Off to the races...









Input and Output

- **InputFormat:**

- `TextInputFormat`
- `KeyValueTextInputFormat`
- `SequenceFileInputFormat`
- ...

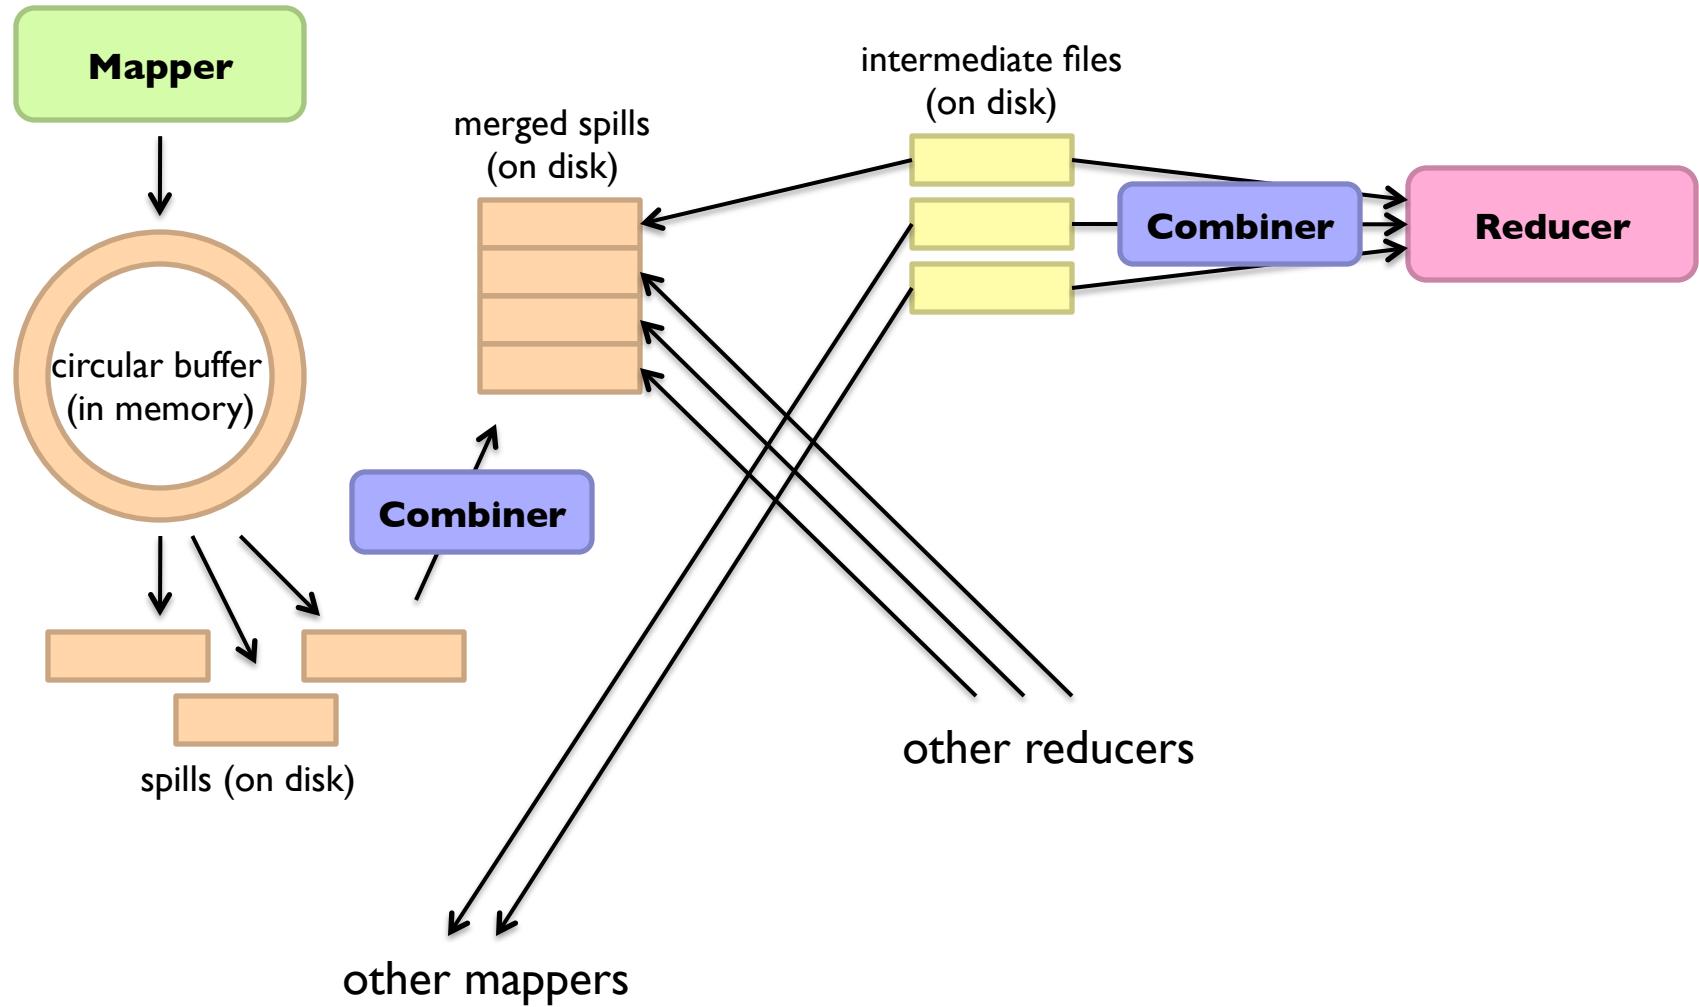
- **OutputFormat:**

- `TextOutputFormat`
- `SequenceFileOutputFormat`
- ...

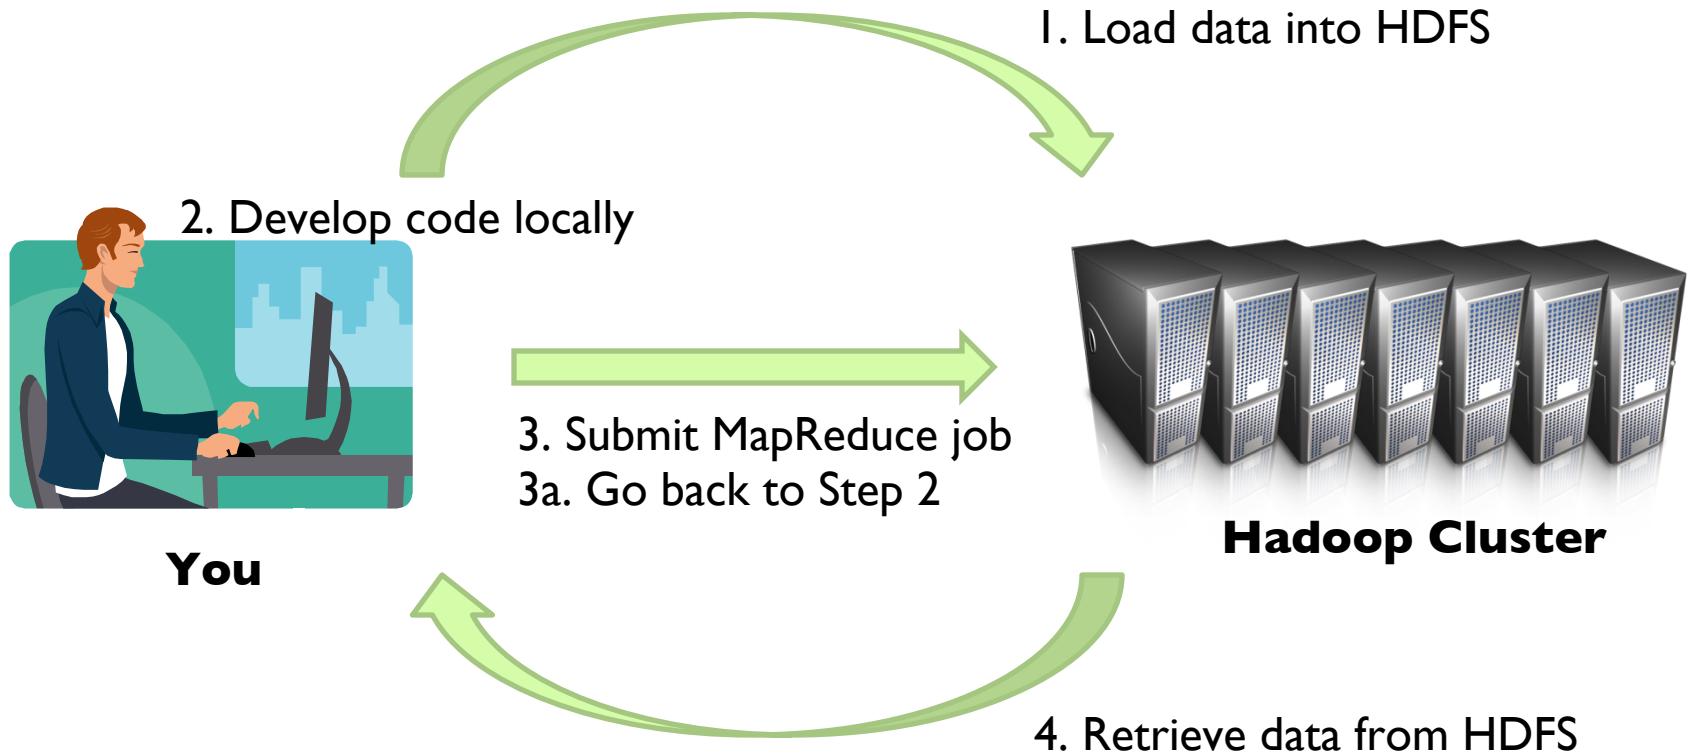
Shuffle and Sort in Hadoop

- Probably the most complex aspect of MapReduce
- Map side
 - Map outputs are buffered in memory in a circular buffer
 - When buffer reaches threshold, contents are “spilled” to disk
 - Spills merged in a single, partitioned file (sorted within each partition): combiner runs during the merges
- Reduce side
 - First, map outputs are copied over to reducer machine
 - “Sort” is a multi-pass merge of map outputs (happens in memory and on disk): combiner runs during the merges
 - Final merge pass goes directly into reducer

Shuffle and Sort



Hadoop Workflow



Debugging Hadoop

- First, take a deep breath
- Start small, start locally
- Build incrementally

Hadoop Debugging Strategies

- Good ol' System.out.println
 - Learn to use the webapp to access logs
 - Logging preferred over System.out.println
 - Be careful how much you log!
- Fail on success
 - Throw RuntimeExceptions and capture state
- Programming is still programming
 - Use Hadoop as the “glue”
 - Implement core functionality outside mappers and reducers
 - Independently test (e.g., unit testing)
 - Compose (tested) components in mappers and reducers

A photograph of a traditional Japanese rock garden. In the foreground, a gravel path is raked into fine, parallel lines. Several large, dark, irregular stones are scattered across the garden. A small, shallow pond is visible in the middle ground, surrounded by more stones and low-lying green plants. In the background, there are more trees and shrubs, and the wooden buildings of a residence are visible behind the garden wall.

Questions?